



UNIVERSITY OF BUCHAREST

FACULTY OF
MATHEMATICS AND
COMPUTER SCIENCE



SPECIALIZATION ARTIFICIAL INTELLIGENCE

Master Thesis

A POST-HOC INTERPRETABILITY ANALYSIS OF XLSTM ARCHITECTURES

GRADUATE

Tănase Victor-Flavian

SCIENTIFIC COORDINATOR

Cristian Rusu PhD

Bucharest, July 2026



UNIVERSITATEA DIN
BUCUREȘTI

FACULTATEA DE
MATEMATICĂ ȘI
INFORMATICĂ



SPECIALIZAREA INTELIGENȚĂ ARTIFICIALĂ

Lucrare de Disertație

UN STUDIU DE INTERPRETABILITATE POST-HOC PENTRU ARHITECTURILE DE TIP XLSTM

Absolvent

Tănase Victor-Flavian

Coordonator științific

Dr Cristian Rusu

București, iulie 2026

Abstract

Recent progress in natural language processing has been largely driven by large-scale neural networks. These architectures aim to learn representations of text that capture syntax, semantics, and long-range dependencies in an efficient and scalable way. Despite significant advances in performance, the internal mechanisms by which these models store and manipulate information remain only partially understood.

This limitation becomes especially relevant in modern architectures that rely on fundamentally different approaches to information storage. Transformer-based models process sequences using self-attention, where information is mixed across tokens at each layer without an explicit notion of persistent state. In contrast, xLSTM introduced by Beck et al. [3] a modern recurrent neural network architecture represents an alternative to the transformer based architectures by incorporating explicit memory that is updated incrementally over time. This allows information to be stored and refined accross consecutive time steps.

This work investigates the internal behavior of the xLSTM architecture through a set of interpretability experiments aimed at understanding how information is represented, stored, and transformed inside the model. The analysis explores multiple levels of abstraction, from token-level dynamics to layer-wise representation structure. Specifically, we apply logit lens[17] [2] to observe the relationship between the internal representations in the embedding space and their mapping into the vocabulary space. Moreover, we compare the hidden states of xLSTM and a comparable Transformer-based model (Mistral 7B) [14] to assess how similarly semantic representations evolve in the two models. Furthermore, we quantify memory behavior by tracking the evolution of the cell state over time and tracking input and forget gate activations at the token level to understand how the model selectively retains or discards information. Together, these experiments provide a multi-perspective view into the internal mechanisms of modern recurrent neural networks, offering interpretable details about how structured memory evolves in large-scale

recurrent architectures.

Finally, we investigate a practical application of xLSTM’s fixed-size memory state by caching contextual information for question-answering tasks. This approach enables inference speedups that increase with context length, as the cached memory can be reused without reprocessing the entire input sequence. We evaluate the method against Mistral 7B and analyze the trade-offs between recurrent and Transformer-based architectures with respect to inference speed and task performance.

Sinopsis

Progresul recent în prelucrarea limbajului natural a fost accelerată de dezvoltarea modelelor lingvistice de mari dimensiuni (LLM). Acest tip de arhitectură are ca scop învățarea unor reprezentări ale limbajului care pot captura proprietăți sintactice, semantice și să modeleze dependențe între cuvinte îndepărtate într-un mod eficient și scalabil. În ciuda îmbunătățirii performanței acestor modele, modul în care ele învață și manipulează informația este doar parțial înțeles.

Această limitare este cu atât mai relevantă în arhitecturile moderne care se bazează pe abordări fundamentale diferite de a reprezenta informația. Modelele de tip transformer procesează secvențe de text cu ajutorul mecanismului de atenție, care folosește informația din tokeni în fiecare strat intern, fără o structură definită pentru memorie. Pe de altă parte, xLSTM o arhitectură de tip recurentă [3] reprezintă o alternativă la modelele de tip transformer, folosind o structură pentru memorie explicită, care este actualizată incremental pe parcursul procesării textului. Acest mecanism are ca avantaj înmagazinarea și rafinarea informației de fiecare dată când un nou token este procesat.

Această lucrare investighează comportamentul arhitecturilor xLSTM printr-o serie de experimente care au ca scop să îmbunătățească modul în care sunt interpretate reprezentarea, înmagazinarea și transformările datelor în interiorul modelului. Analiza se desfășoară pe mai multe nivele de abstractizare, începând de la analiza la nivel de token, ca mai apoi să fie observate reprezentările la nivel de strat. Mai specific, sunt utilizate tehnici de tip logit lens [17] [2] pentru a remarca evoluția relației dintre reprezentările intermediare în spațiul de embeddings și spațiul vocabularului. De asemenea, comparăm reprezentările interne ale modelului xLSTM cu cele ale unei arhitecturi comparabile de tip transformer, anume Mistral-7B [14]. Scopul comparației este de a observa similaritatea evoluției reprezentărilor semantice dintre două arhitecturi diferite. Totodată, analizăm modificarea memoriei pe parcursul procesării tokenilor, urmărind evoluția memoriei pe parcursul parcurgerii unui text. și urmărim activările porților input și forget la nivel de token pentru a înțelege cum

modelul alege să rețină și să elimine selectiv conext. Împreună aceste experimente oferă o perspectivă nuanțată a interpretării mecanismelor din interiorul arhitecturii xLSTM.

În final, explorăm o utilizare practică a stării de memorie de dimensiune fixă din xLSTM prin memorarea informației contextuale în scenarii care testează abilitatea modelului de a răspunde la întrebări. Această abordare permite îmbunătățirea vitezei de răspuns pe măsură ce dimensiunea contextului crește, deoarece starea de memorie stocată poate fi reutilizată fără a procesa întreaga secvență de intrare. Această metodă cu modelul Mistral 7B și analizăm avantajele și dezavantajele arhitecturilor recurente și celor bazate pe Transformere din perspectiva vitezei și a performanței de fiecare dată.

Contents

1	Introduction	8
1.1	Problem	8
1.2	Motivation	9
1.3	Contribution	10
1.4	Outline	11
2	Preliminaries and Related Work	12
2.1	Recurrent Neural Networks	12
2.2	Long-Short Term Memory	14
2.3	Extended Long-Short Term Memory	16
2.3.1	Addressing Classic LSTM Limitations	16
2.3.2	sLSTM	17
2.3.3	mLSTM	18
2.3.4	The xLSTM Block	19
3	Method	21
3.1	Research Objectives	21
3.2	Models	21
3.2.1	xLSTM-7B	21
3.2.2	Mistral-7B	23
3.3	Baseline evaluation	24
3.4	Interpretability Study	26
3.4.1	The Logit Lens	27
3.4.2	Hidden State Similarity with CKA	28
4	Experiments	29
5	Conclusion and Future Work	30

1 Introduction

Since the inception of the Transformer [28], large architectures utilizing stacked attention mechanisms have dominated the research landscape, exhibiting impressive performance in natural language processing tasks. This architectural paradigm scaled rapidly, progressing from hundreds of millions of parameters to billions, and ultimately to trillions for enterprise-scale models. In parallel, with this scaling trend, a significant part of the scientific community focused their efforts on understanding the internal mechanisms of these models. However, this intense scientific work has focused almost exclusively towards attention-based architectures, leaving non-transformer based models comparatively less explored.

Despite their success, Transformers suffer from a fundamental computational bottleneck: the self-attention mechanism scales quadratically ($\mathcal{O}(N^2)$) with respect to sequence length, making long-context processing computationally expensive. This limitation has led to a recent resurgence of interest in recurrent architectures and state-based architectures that aim to achieve linear-time efficiency ($\mathcal{O}(N)$) without sacrificing performance. Key contributions in this direction include: Structured State Space Models (SSMs) like Mamba [11], linear transformers such as Receptance Weighted Key Value (RWKV)[21] and Extended Long-Short Term Memory(xLSTM) [3]. The xLSTM architecture extends the classic LSTM cell [13] by introducing exponential gating and matrix-valued memory states with gated updates. These changes look to address the scaling problem of the Transformers and the storage bottlenecks of traditional LSTMs, achieving comparable performance while maintaining linear-time inference.

1.1 Problem

Deep architectures have produced tremendous advancements in the field of machine learning over the past decade. Their pattern recognition and information representation capabilities are yet to be completely understood even to this day. Post-hoc interpretability

in deep neural networks is a challenging problem due to several intrinsic mathematical and structural properties of these models. First, modern architectures operate in high-dimensional representation spaces, where features are typically distributed across multiple dimensions rather than localized to individual neurons. Moreover, multiple concepts may be encoded within overlapping directions of the representation space. Second, neural networks apply a sequence of non-linear transformations, making the relationship between input tokens and internal representations difficult to describe using simple analytical tools. As network depth increases, interactions between representations become progressively more complex. Finally, although the computation of hidden states is fully specified by the model parameters and architecture, the resulting computation graph is generally too complex to provide specific human-interpretable explanations of how particular representations or predictions are produced.

1.2 Motivation

The motivation for this work originates from the recent resurgence of recurrent architectures in large-scale sequence modelling. While analyzing the current state of research on recurrent neural networks, the xLSTM architecture, which Beck et al.[3] proposed in 2024, became a personal point of interest. As one of the first recurrent architectures in recent years to demonstrate competitive performance at scales traditionally dominated by Transformer-based models, xLSTM represents an important development in the ongoing search for efficient alternatives to self-attention.

While the performance of Transformer-based language models has been extensively studied, a substantial body of interpretability research has also emerged to investigate their internal mechanisms. In contrast, the behavior of modern recurrent architectures such as xLSTM remains comparatively underexplored. This observation motivated the present work, which seeks to apply and adapt interpretability techniques inspired by previous studies of neural language models, including the Logit Lens methodology and analyses of recurrent network dynamics.

By investigating how information is represented, stored, and updated within xLSTM, this thesis aims to contribute to a better understanding of the model’s memory mechanisms and internal representations. Such insights may help validate architectural design choices, identify limitations, and inform future developments of memory-based language models.

Another question to be answered is if an alternative option to the transformer models like the xLSTM architecture can operate in a practical environment. In parallel, it remains to be understood how competing architectures can displace Transformers, which are widely adopted in modern infrastructure through highly optimized hardware acceleration, making them exceptionally difficult to replace in production workflows.

1.3 Contribution

To address the lack of interpretability data on scalable recurrent architectures, the primary objective of this thesis is to execute a post-hoc evaluation framework on the internal representations and state dynamics of the xLSTM-7B architecture. This work maps the internal mechanics of the model by moving systematically from macro-level validation down to micro-level gating operations.

The first phase of the research establishes an empirical baseline by evaluating the pre-trained xLSTM-7B model across standard downstream zero-shot datasets utilized in the foundational literature, specifically LAMBADA (OpenAI)[19], PIQA[5], Winogrande[24], ARC-Challenge[8], ARC-Easy[8], and Hellaswag[31]. This behavioral evaluation serves to validate the functional performance limits of the model prior to any internal analysis. Following this initial baseline, the second phase adapts the logit lens[17] [2] technique to project intermediate hidden states from the recurrent layers directly onto the vocabulary space. Initial observations from this tracking indicate that vocabulary prediction convergence across intermediate layers follows a trajectory structurally similar to that observed

in standard Transformer models.

The third phase compares how hidden representations evolve across the layers of a recurrent network compared to a Transformer, using xLSTM-7B and Mistral-7B. The final two phases of the interpretability study focus on the temporal dynamics of the architecture. Specifically, they examine the evolution of the memory state over time and the contribution of individual tokens to long-term memory through the input and forget gates. To analyze these mechanisms, we process long token sequences and track the Frobenius norm of the cell state, together with the activations of the input and forget gates. This allows us to observe how information is stored, retained, and updated throughout the sequence.

On top of the interpretability experiments, we evaluate the practical utility of the xLSTM model in a question answering setting. We leverage the constant-size memory state of the model and explore caching it when processing context that is later reused for answering questions. This approach enables an inference speedup that increases with the size of the context, as the stored memory can be reused without recomputing the full sequence.

1.4 Outline

2 Preliminaries and Related Work

To provide the necessary background for the evaluation of the xLSTM-7B architecture, this chapter presents the theoretical concepts and related work relevant to this study. The first section reviews the evolution of recurrent neural networks, from classical RNNs to the modern architecture proposed by Beck et al. [3] that incorporate gating mechanisms and matrix-based memory structures.

2.1 Recurrent Neural Networks

Recurrent Neural Networks (RNNs) represent one of the earliest neural architectures specifically designed for sequence modeling tasks. Their foundations can be traced to the work of Rumelhart et al. [23], who introduced the backpropagation algorithm and discussed neural networks with recurrent connections. Unlike feedforward neural networks, which process each input independently, RNNs maintain a hidden state that evolves over time, allowing information from previous inputs to influence future predictions.

The key idea behind recurrent architectures is the reuse of the same set of parameters at every time step. Given an input sequence, the hidden state is updated recursively based on both the current input and the previous hidden state, enabling the network to capture temporal dependencies within the data. This property made RNNs a natural choice for tasks involving sequential information, such as language modeling, speech recognition, machine translation, and time-series forecasting.

The operation of a standard RNN can be formally described through its hidden state update (equation 1). At a given timestamp $t = \overline{1, T}$, for an input x_t , the hidden state h_t is computed as:

$$h_t = \phi(W_{xh}x_t + W_{hh}h_{t-1} + b_h), \quad (1)$$

where W_{xh} and W_{hh} are learnable weight matrices, b_h is a bias term, and $\phi(\cdot)$ denotes




Figure 1: RNN: compressed (left) and unfolded(right), diagram inspired by [18]

a non-linear activation function such as tanh or sigmoid. The hidden state is then used to produce the output y_t :

$$y_t = W_{hy}h_t + b_y, \quad (2)$$

where W_{hy} and b_y are the output weight matrix and bias, respectively. Through the recurrent connection involving h_{t-1} , the network is able to propagate information across time steps, allowing previous inputs to influence future predictions.

Training recurrent neural networks requires propagating gradients not only through the layers of the network but also through time. This is typically achieved using *Backpropagation Through Time* (BPTT), introduced by Werbos et al.[29]. BPTT unfolds the recurrent network across the input sequence, transforming it into an equivalent deep feedforward network with shared parameters at each time step. Gradients are then computed by applying the backpropagation algorithm to the unfolded computational graph and accumulated across all time steps before updating the model parameters.

While this training procedure allows RNNs to learn temporal dependencies, it also exposes a major limitation of the architecture. During BPTT, gradients are propagated backward through all previous time steps of the sequence. At each step, the gradient is repeatedly multiplied by the same recurrent transformation, as it is passed through the unfolded network in time. As a result, the gradient magnitude may gradually decrease

or increase with each multiplication. When it becomes too small, the network struggles to learn from earlier time steps, a phenomenon known as the *vanishing gradient problem*. Conversely, when it grows too large, training becomes unstable, leading to the *exploding gradient problem* [4].

These issues significantly limit the ability of standard RNNs to capture long-range dependencies in sequential data. To address the exploding gradient problem, Pascanu et al. [20] introduced *gradient clipping*, a technique that limits the magnitude of gradients during training and improves optimization stability. Nevertheless, the difficulty of preserving information over long sequences remained a fundamental challenge, motivating the development of more advanced recurrent architectures such as Long Short-Term Memory (LSTM)[13] networks and Gated Recurrent Units (GRUs)[6][7].

2.2 Long-Short Term Memory

The **Long Short-Term Memory (LSTM)** architecture, introduced by Hochreiter and Schmidhuber [13] in the 1990s, was proposed to address the limitations of standard RNNs when learning long-range dependencies. The key innovation of the architecture is the introduction of a dedicated memory state that can preserve information across many time steps. Instead of relying solely on the hidden state, the LSTM uses a set of gating mechanisms to control how information is written to, stored in, and retrieved from memory. This allows the model to memorize relevant information over longer sequences while reducing the impact of the vanishing gradient problem encountered in traditional recurrent networks.

The standard LSTM cell maintains an internal cell state, $c_t(6)$, which acts as a long-term memory buffer, and a hidden state, $h_t(8)$, which serves as the short-term working memory routed to subsequent layers. The flow of information inside the cell is controlled by three gates with different scopes. The forget gate(3) manages what proportion of the historical




Figure 2: The LSTM cell, figure inspired by [18]

cell state to discard by analyzing the current input and the previous hidden state. The input gate(4) controls how much new information, computed from the current input and previous hidden state, is incorporated into the cell state. Finally, the output gate(7) determines which information stored in the updated cell state is propagated as the next hidden state.

$$f_t = \sigma(W_f x_t + R_f h_{t-1} + b_f) \quad (3)$$

$$i_t = \sigma(W_i x_t + R_i h_{t-1} + b_i) \quad (4)$$

$$\tilde{c}_t = \tanh(W_c x_t + R_c h_{t-1} + b_c) \quad (5)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t \quad (6)$$

$$o_t = \sigma(W_o x_t + R_o h_{t-1} + b_o) \quad (7)$$

$$h_t = o_t \odot \tanh(c_t) \quad (8)$$

Overall, the introduction of a dedicated memory cell and gating mechanisms allows

LSTMs to address the limitations of standard RNNs in capturing long-range dependencies. However, the sequential processing over the time dimension still limits parallelization and makes them computationally expensive for long sequences. Several architectural extensions were later proposed to improve their effectiveness in different settings. Bidirectional recurrent networks [25] improved contextual understanding by processing sequences in both forward and backward directions, while encoder-decoder architectures [27] enabled sequence-to-sequence learning by separating the encoding and decoding processes. In addition, stacked recurrent architectures [10] increased model capacity by composing multiple recurrent layers, allowing for more expressive hierarchical representations.

Despite these improvements, these variants did not remove the fundamental sequential bottleneck of recurrent computation, which limited scalability on long sequences. These limitations motivated the transition toward architectures that avoid recurrence entirely, most notably Transformer[28] models based on self-attention. In this broader context, recent work has revisited recurrent memory mechanisms at scale, leading to architectures such as xLSTM, which aim to combine the efficiency of structured recurrence with improved scalability and expressive memory modeling.

2.3 Extended Long-Short Term Memory

The extended LSTM (xLSTM) architecture was introduced by Beck et al. [3] to overcome these bottlenecks by leveraging modern techniques while fundamentally modifying the LSTM memory structure.

2.3.1 Addressing Classic LSTM Limitations

To understand the architectural shifts in the xLSTM, it is necessary to outline the three main bottlenecks of the classic LSTM:

1. **Inability to revise storage decisions:** Traditional LSTMs struggle to update or

revise a stored value when new, more relevant information is found.

2. **Limited storage capacities:** Information in a classic LSTM is compressed into a scalar cell state, leading to poor performance when predicting rare tokens or retaining extensive context.
3. **Lack of parallelizability:** The recurrent memory update relies on hidden-to-hidden connections between consecutive time steps, requiring strictly sequential processing and preventing the highly parallel computation that characterizes Transformer architectures.

To address these issues, xLSTM introduces several changes to the classical LSTM architecture. First, it replaces the traditional logistic gating mechanism with *exponential gating*, which provides a less restricted range for controlling memory updates. To ensure numerical stability during training and inference, xLSTM uses normalization and stabilization techniques. In addition, the architecture looks to improve the design of the memory state, resulting in two distinct recurrent building blocks: the *scalar LSTM* (*sLSTM*) and the *matrix LSTM* (*mLSTM*). While sLSTM extends the traditional LSTM memory formulation with improved gating and memory management, mLSTM introduces a matrix-based memory structure capable of storing richer associations between tokens.

2.3.2 sLSTM

The sLSTM (scalar LSTM) remains structurally closer to the classic architecture but introduces exponential activation functions for the input and forget gates. This modification allows the model to control its storage more effectively:

$$\mathbf{c}_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + \mathbf{i}_t \odot \mathbf{z}_t \quad \text{cell state} \quad (9)$$

$$\mathbf{h}_t = \mathbf{o}_t \odot \tilde{\mathbf{h}}_t, \quad \tilde{\mathbf{h}}_t = \psi(\mathbf{c}_t) \quad \text{hidden state} \quad (10)$$

$$\mathbf{z}_t = \varphi(\tilde{\mathbf{z}}_t), \quad \tilde{\mathbf{z}}_t = W_z x_t + R_z h_{t-1} + b_z \quad \text{cell input} \quad (11)$$

$$\mathbf{i}_t = \sigma(\tilde{\mathbf{i}}_t), \quad \tilde{\mathbf{i}}_t = W_i x_t + R_i h_{t-1} + b_i \quad \text{input gate} \quad (12)$$

$$\mathbf{f}_t = \sigma(\tilde{\mathbf{f}}_t), \quad \tilde{\mathbf{f}}_t = W_f x_t + R_f h_{t-1} + b_f \quad \text{forget gate} \quad (13)$$

$$\mathbf{o}_t = \sigma(\tilde{\mathbf{o}}_t), \quad \tilde{\mathbf{o}}_t = W_o x_t + R_o h_{t-1} + b_o \quad \text{output gate} \quad (14)$$

To prevent the numerical overflows specific to the exponential function, a normalizer state (n_t) and stabilization states are utilised:

$$m_t = \max(\log(f_t) + m_{t-1}, \log(i_t)) \quad \text{stabilizer state} \quad (15)$$

$$i'_t = \exp(\log(i_t) - m_t) = \exp(\tilde{i}_t - m_t) \quad \text{stabil. input gate} \quad (16)$$

$$f'_t = \exp(\log(f_t) + m_{t-1} - m_t) \quad \text{stabil. forget gate} \quad (17)$$

Despite these changes, because it retains hidden-hidden recurrent connections, the sLSTM doesn't improve the sequential processing bottleneck of the classic LSTM. For the scope of this work, the architecture used in the experiments solely relies on mLSTM cells.

2.3.3 mLSTM

The mLSTM fundamentally shifts the LSTM paradigm to solve both the storage capacity and parallelization bottlenecks. To overcome the limited capacity of scalar states, the mLSTM expands the memory cell into a matrix $C \in \mathbb{R}^{d \times d}$ where d represents the embedding dimension of x_t . Crucially, to achieve this efficiency, the mLSTM entirely abandons memory mixing (the hidden-hidden recurrent connections) allowing parallelization.

Relying on Bidirectional Associative Memory (BAM)[16] principles and Transformer terminology, the mLSTM stores information as key-value pairs (k_t, v_t) using a covariance

(outer product) update rule, and retrieves it using a query vector (q_t). The mLSTM cell also uses the stabilizer state introduced in the sLSTM equations(15)

$$C_t = f_t C_{t-1} + i_t v_t k_t^\top \quad \text{cell state} \quad (18)$$

$$n_t = f_t n_{t-1} + i_t k_t \quad \text{norm. state} \quad (19)$$

$$h_t = \mathbf{o}_t \odot \tilde{h}_t, \quad \tilde{h}_t = C_t q_t / \max \{|n_t^\top q_t|, 1\} \quad \text{hidden state} \quad (20)$$

$$q_t = W_q x_t + b_q \quad \text{query input} \quad (21)$$

$$k_t = \frac{1}{\sqrt{d}} W_k x_t + b_k \quad \text{key input} \quad (22)$$

$$v_t = W_v x_t + b_v \quad \text{value input} \quad (23)$$

$$i'_t = \exp(\tilde{i}_t - m_t), \quad \tilde{i}_t = w_i^\top x_t + b_i \quad \text{input gate} \quad (24)$$

$$f'_t = \exp(\log(\sigma(\tilde{f}_t)) + m_{t-1} - m_t), \quad \tilde{f}_t = w_f^\top x_t + b_f \quad \text{forget gate} \quad (25)$$

$$\mathbf{o}_t = \sigma(\tilde{\mathbf{o}}_t), \quad \tilde{\mathbf{o}}_t = W_o x_t + b_o \quad \text{output gate} \quad (26)$$

2.3.4 The xLSTM Block

The xLSTM block leverages the mLSTM as a multi-head sequence-mixing mechanism. First, the input token representation (x_t) is normalized using RMSNorm [32]. The normalized vector is then linearly projected into multiple heads. Each head receives its own set of projected query, key, value, and gate vectors and processes them independently through a dedicated mLSTM cell. As a result, the recurrent state is also maintained separately for each head. The outputs of the individual heads, denoted as ($h_1, h_2, \dots, h_N$), are then normalized per head, gated, concatenated, and passed through a shared linear projection layer. This produces an output vector with the same dimensionality as the input, allowing it to be added back via a residual connection. After the mLSTM sub-layer, the block applies a second pre-RMSNorm residual sub-layer: a SwiGLU feed-forward network. This layer does not mix information across time. Instead, it processes each token representation independently. The final output of the block is obtained by adding the FFN output back to the result of the mLSTM residual branch.

This block structure is the one used in the available pre-trained xLSTM-7b model.




Figure 3: mLSTM Block

It differs from the original xLSTM paper, where the mLSTM block used a more specialized structure with an up-projection, causal convolution, block-diagonal query/key projections, and an internal skip connection. In the pre-trained model, the block is closer to a standard transformer block: a multi-head mLSTM sequence-mixing layer followed by a SwiGLU feed-forward layer. For the scope of this study, we'll refer to this version of the mLSTM block.

3 Method

This section focuses on the technical details of the experiments that are presented in this work.

3.1 Research Objectives

This work has three clear goals that can be defined as follows:

1. **Baseline evaluation of the xLSTM-7B model**(3.3): asses the performance of the model by evaluating it and compare it to a comparable transformer model, the Mistral-7B
2. **Interpretability study**: dissect the xLSTM components and investigate the internal mechanisms responsible for memory formation, information retention, and representation development.
3. **Practical Usage and generation speedup: Efficient Inference**: evaluate whether the fixed-size memory state of xLSTM can be leveraged to reduce inference time while maintaining task performance in practical setup.

3.2 Models

3.2.1 xLSTM-7B

In this study, we use the pretrained xLSTM-7B model available on Hugging Face. As discussed in the previous section, the implementation differs slightly from the architecture described in the original xLSTM paper, primarily through a modified version of the mLSTM block (Figure 3). The model follows a decoder-only language modeling architecture. Input tokens are first mapped to dense vector representations through an embedding layer and subsequently processed by a stack of 32 mLSTM blocks. The resulting hidden representations are then normalized using Root Mean Square Normalization (RMSNorm) before being projected into the vocabulary space through a linear output layer. Figure 4 provides an overview of the complete model architecture.




Figure 4: xLSTM: architectural overview, figure after

The model contains 6,865,424,896 parameters and uses an embedding dimension of 4096. The hidden state dimension is also 4096 and is split across 8 parallel heads, each operating on a reduced subspace of dimension $(4096 / 8 = 512)$. Due to the recurrent nature of the architecture, the model can in principle process sequences of arbitrary length, as its memory is maintained through the hidden and cell states rather than being constrained by a fixed positional context window. The model is used in its pretrained form, and no

additional training or fine-tuning is performed in this study.

3.2.2 Mistral-7B

For comparison purposes, we use the pretrained Mistral-7B-v0.3 model as a representative Transformer-based architecture. Like xLSTM, Mistral is a decoder-only language model and is used in its pretrained form without additional fine-tuning. However, the mechanism by which information is processed differs fundamentally between the two architectures.

Mistral relies on Grouped-Query Attention (GQA)[1] as its primary sequence modeling mechanism, utilizing 32 query heads and 8 key-value heads to optimize KV cache memory efficiency. The architecture uses Rotary Positional Embeddings (RoPE) [26] to natively sustain a 32,768-token context window. Within each Transformer decoder block, the embedded tokens are processed using RMS normalization before entering the GQA layer, where separate linear projections generate the query, key, and value vectors. The attention outputs are combined, passed through a residual connection, and processed by a subsequent RMS normalization layer. This is followed by a feed-forward network consisting of three linear transformations featuring a gated SiLU activation mechanism (SwiGLU). A final residual connection completes the block.

The Mistral-7B architecture consists of 32 Transformer decoder layers, each containing a multi-head self-attention module and a feed-forward network. The model uses a hidden dimension of 4096 and approximately 7 billion parameters, making it comparable in scale to the xLSTM-7B model. This architectural similarity in size allows differences observed in the experiments to be attributed primarily to the underlying sequence modeling mechanisms rather than to disparities in model capacity.

Throughout this study, Mistral serves as a reference model for both performance benchmarking and representation analysis, enabling a comparison between attention-based and recurrent approaches to language modeling.




Figure 5: Mistral-7B-v0.3: architectural overview

3.3 Baseline evaluation

To establish a baseline understanding of the model’s capabilities, we evaluate xLSTM-7B on a collection of widely used language understanding and reasoning benchmarks. These datasets assess different aspects of language modeling, ranging from commonsense reasoning and physical intuition to reading comprehension and scientific knowledge.

- **HellaSwag** [31]: A commonsense reasoning benchmark consisting of partially described situations followed by multiple candidate continuations. The task requires

selecting the most plausible ending among several alternatives. HellaSwag is designed to be challenging for language models by including distractor options generated through adversarial filtering, making it a strong test of contextual understanding and commonsense reasoning.

- **LAMBADA** [19]: A language modeling benchmark that evaluates the ability to predict the final word of a passage using the broader context of the document. Success on this task requires integrating information across multiple sentences and is therefore particularly relevant for assessing long-range dependency modeling.
- **ARC-Challenge** [8]: The challenging subset of the AI2 Reasoning Challenge dataset. It consists of multiple-choice science questions that are difficult to answer using simple retrieval or pattern matching techniques. The benchmark evaluates reasoning ability, scientific knowledge, and the capacity to combine information from multiple sources.
- **ARC-Easy** [8]: A simpler subset of the ARC benchmark containing elementary-level science questions. While less demanding than ARC-Challenge, it provides a useful measure of basic factual knowledge and reasoning capabilities.
- **PIQA** [5]: The Physical Interaction Question Answering benchmark evaluates physical commonsense reasoning. Questions describe everyday situations and require selecting the most plausible solution to a practical problem. The benchmark focuses on a model’s understanding of how objects and actions interact in the physical world.
- **WinoGrande** [24]: A large-scale commonsense reasoning benchmark derived from the Winograd Schema Challenge. The task involves resolving ambiguous references in a sentence by selecting the correct antecedent. Performance on WinoGrande is often used as a measure of a model’s ability to perform contextual and commonsense reasoning beyond simple pattern recognition.

Together, these benchmarks provide a broad evaluation of the model across several dimen-

sions, including commonsense reasoning, long-context understanding, physical reasoning, scientific knowledge, and language comprehension.

3.4 Interpretability Study

The study of interpretable properties in deep neural networks has helped researchers better understand the internal mechanisms of otherwise opaque models. In computer vision, Zeiler et al. [30] improved the understanding of convolutional neural networks by using a deconvolutional framework to project intermediate feature activations back into the input space. This made it possible to observe how increasingly abstract visual representations emerge throughout the network layers.

In recurrent and Transformer-based architectures, interpretability research has focused on understanding what information is stored in hidden representations. Hewitt and Manning [12] showed that the hidden states of both RNNs and Transformers encode syntactic structure. Using a structural probing framework, they demonstrated that syntax trees can be recovered through a learned linear transformation of the representation space. Their results on models such as ELMo [22] and BERT [9] suggest that linguistic structure is organized within the geometry of the learned representations rather than being stored as unstructured information.

This interpretability study follows a layered approach that moves from high-level model outputs to internal mechanisms. We first analyze intermediate representations to understand what information is available at different depths of the network. We then compare these representations across architectures to study structural similarities and differences. Finally, we focus on the internal memory dynamics of xLSTM, examining how information is stored and controlled through its recurrent state and gating mechanisms. This progressive design allows us to move from observable predictions to the underlying processes that generate them.

3.4.1 The Logit Lens

The logit lens[17][2] is an experimental framework which studies the evolution of the latent representations in the network. To track how predictions form as text moves through the architecture, we extracted the hidden state vectors from each internal block.

Under normal circumstances, word predictions are made at the very end of the network. In the logit lens setup, at each intermediate block, we take the temporary hidden states, apply the normalization then apply the learned linear(the language modeling head) projection into the vocabulary space.

To evaluate how these internal guesses change and improve across depth, we analyzed the resulting word distributions in three ways:

- **Logits (Top-1 Predictions):** We evaluated the raw output scores across the entire vocabulary to identify which single token was the model’s top choice at each individual block. This allowed us to inspect the literal words the network considers during mid-network calculation.
- **Target Ranks:** Even when an early layer does not select the final correct token as its top choice, that token may still be highly rated. We tracked the strict rank order of the final true token within each block’s internal probability distribution. Monitoring how this rank climbs toward the first position shows how quickly the network narrows down its uncertainty.
- **Kullback-Leibler (KL) Divergence:** Computing the KL divergence, helped measuring the overall statistical difference between the token probability distribution of an intermediate block and that of the final output block. This provided a continuous metric to monitor how smoothly the network’s internal beliefs converge toward its final output, and to verify how quickly the original input representation is transformed upon entering the network.

3.4.2 Hidden State Similarity with CKA

Building on the logit lens analysis, which was originally introduced for Transformer-based architectures, we next investigate how the representations learned by a recurrent architecture such as xLSTM compare to those of a Transformer model, specifically Mistral-7B. The transformer model choice was driven by the fact that both models share the same embedding dimension and the same number of layers, which allows for a more aligned layer-wise analysis.

To evaluate differences in structural representations between the two architectures, we compare the hidden representations of each model using Centered Kernel Alignment (CKA) [15], a metric designed for measuring similarity between neural network representations. In this study, we use linear CKA computed over mean-centered hidden states across tokens, which quantifies similarity through normalized Frobenius norms of cross-covariance matrices:

$$X_c = X - \frac{1}{B} \sum_{i=1}^B X_{i,\cdot} \quad (27)$$

$$Y_c = Y - \frac{1}{B} \sum_{i=1}^B Y_{i,\cdot} \quad (28)$$

where B represents the batch dimension

$$\text{Linear CKA}(X, Y) = \frac{\|Y_c^\top X_c\|_F^2}{\|X_c^\top X_c\|_F \|Y_c^\top Y_c\|_F} \quad (29)$$

For this purpose, we extract hidden state activations from each block of both models using the same input prompt, ensuring a consistent basis for comparison.

4 Experiments

5 Conclusion and Future Work

References

- [1] Joshua Ainslie, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebrón, and Sumit Sanghai. *GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints*. 2023. arXiv: [2305.13245](https://arxiv.org/abs/2305.13245) [cs.CL]. URL: <https://arxiv.org/abs/2305.13245>.
- [2] Belrose et al. “Eliciting Latent Predictions from Transformers with the Tuned Lens”. In: *arXiv preprint arXiv:2303.08112* (2023).
- [3] Maximilian Beck, Korbinian Pöppel, Markus Spanring, Andreas Auer, Oleksandra Prudnikova, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. *xLSTM: Extended Long Short-Term Memory*. 2024. arXiv: [2405.04517](https://arxiv.org/abs/2405.04517) [cs.LG]. URL: <https://arxiv.org/abs/2405.04517>.
- [4] Yoshua Bengio, Patrice Simard, and Paolo Frasconi. “Learning Long-Term Dependencies with Gradient Descent is Difficult”. In: *IEEE Transactions on Neural Networks* (1994).
- [5] Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. “Piqa: Reasoning about physical commonsense in natural language”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 34. 05. 2020, pp. 7432–7439.
- [6] Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. *Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation*. 2014. arXiv: [1406.1078](https://arxiv.org/abs/1406.1078) [cs.CL]. URL: <https://arxiv.org/abs/1406.1078>.
- [7] Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. *Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling*. 2014. arXiv: [1412.3555](https://arxiv.org/abs/1412.3555) [cs.NE]. URL: <https://arxiv.org/abs/1412.3555>.
- [8] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. “Think you have solved question answering? try arc, the ai2 reasoning challenge”. In: *arXiv preprint arXiv:1803.05457* (2018).

- [9] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding”. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Ed. by Jill Burstein, Christy Doran, and Thamar Solorio. Minneapolis, Minnesota: Association for Computational Linguistics, June 2019, pp. 4171–4186. DOI: [10.18653/v1/N19-1423](https://doi.org/10.18653/v1/N19-1423). URL: <https://aclanthology.org/N19-1423/>.
- [10] Alex Graves. “Generating Sequences With Recurrent Neural Networks”. In: *arXiv preprint arXiv:1308.0850* (2013).
- [11] Albert Gu and Tri Dao. *Mamba: Linear-Time Sequence Modeling with Selective State Spaces*. 2024. arXiv: [2312.00752](https://arxiv.org/abs/2312.00752) [cs.LG]. URL: <https://arxiv.org/abs/2312.00752>.
- [12] John Hewitt and Christopher D. Manning. “A Structural Probe for Finding Syntax in Word Representations”. In: *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Ed. by Jill Burstein, Christy Doran, and Thamar Solorio. Minneapolis, Minnesota: Association for Computational Linguistics, June 2019, pp. 4129–4138. DOI: [10.18653/v1/N19-1419](https://doi.org/10.18653/v1/N19-1419). URL: <https://aclanthology.org/N19-1419/>.
- [13] Sepp Hochreiter and Jürgen Schmidhuber. “LSTM can solve hard long time lag problems”. In: *Proceedings of the 10th International Conference on Neural Information Processing Systems*. NIPS’96. Denver, Colorado: MIT Press, 1996, pp. 473–479.
- [14] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, Léo Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El

- Sayed. *Mistral 7B*. 2023. arXiv: [2310.06825](https://arxiv.org/abs/2310.06825) [cs.CL]. URL: <https://arxiv.org/abs/2310.06825>.
- [15] Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. *Similarity of Neural Network Representations Revisited*. 2019. arXiv: [1905.00414](https://arxiv.org/abs/1905.00414) [cs.LG]. URL: <https://arxiv.org/abs/1905.00414>.
- [16] Bart Kosko. “Adaptive bidirectional associative memories”. In: *Applied optics* 26.23 (1987), pp. 4947–4960.
- [17] nostalgebraist. *Interpreting GPT: The Logit Lens*. 2020. URL: <https://www.lesswrong.com/posts/AcKRB8wDpdaN6v6ru/interpreting-gpt-the-logit-lens>.
- [18] Christopher Olah. *Understanding LSTM Networks*. 2015. URL: <https://colah.github.io/posts/2015-08-Understanding-LSTMs/> (visited on 06/01/2026).
- [19] Denis Paperno, Germán Kruszewski, Angeliki Lazaridou, Ngoc Quan Pham, Raffaella Bernardi, Sandro Pezzelle, Marco Baroni, Gemma Boleda, and Raquel Fernández. “The LAMBADA dataset: Word prediction requiring a broad discourse context”. In: *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 2016, pp. 1525–1534.
- [20] Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. “On the Difficulty of Training Recurrent Neural Networks”. In: *ICML*. 2013.
- [21] Bo Peng, Eric Alcaide, Quentin Anthony, Alon Albalak, Samuel Arcadinho, Stella Biderman, Huanqi Cao, Xin Cheng, Michael Chung, Matteo Grella, Kranthi Kiran GV, Xuzheng He, Haowen Hou, Jiaju Lin, Przemyslaw Kazienko, Jan Kocon, Jiaming Kong, Bartłomiej Koptyra, Hayden Lau, Krishna Sri Ipsit Mantri, Ferdinand Mom, Atsushi Saito, Guangyu Song, Xiangru Tang, Bolun Wang, Johan S. Wind, Stanislaw Wozniak, Ruichong Zhang, Zhenyuan Zhang, Qihang Zhao, Peng Zhou, Qinghua Zhou, Jian Zhu, and Rui-Jie Zhu. *RWKV: Reinventing RNNs for the Transformer Era*. 2023. arXiv: [2305.13048](https://arxiv.org/abs/2305.13048) [cs.CL]. URL: <https://arxiv.org/abs/2305.13048>.

- [22] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. *Deep contextualized word representations*. 2018. arXiv: [1802.05365](https://arxiv.org/abs/1802.05365) [cs.CL]. URL: <https://arxiv.org/abs/1802.05365>.
- [23] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. “Learning representations by back-propagating errors”. In: *nature* 323.6088 (1986), pp. 533–536.
- [24] Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. “Winogrande: An adversarial winograd schema challenge at scale”. In: *Communications of the ACM* 64.9 (2021), pp. 99–106.
- [25] Mike Schuster and Kuldip K. Paliwal. “Bidirectional Recurrent Neural Networks”. In: *IEEE Transactions on Signal Processing* (1997).
- [26] Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. *RoFormer: Enhanced Transformer with Rotary Position Embedding*. 2023. arXiv: [2104.09864](https://arxiv.org/abs/2104.09864) [cs.CL]. URL: <https://arxiv.org/abs/2104.09864>.
- [27] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. “Sequence to Sequence Learning with Neural Networks”. In: *NeurIPS* (2014).
- [28] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. “Attention is all you need”. In: *Advances in neural information processing systems* 30 (2017).
- [29] Paul J. Werbos. “Backpropagation Through Time: What It Does and How to Do It”. In: *Proceedings of the IEEE* (1990).
- [30] Matthew D Zeiler and Rob Fergus. *Visualizing and Understanding Convolutional Networks*. 2013. arXiv: [1311.2901](https://arxiv.org/abs/1311.2901) [cs.CV]. URL: <https://arxiv.org/abs/1311.2901>.
- [31] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. “HellaSwag: Can a Machine Really Finish Your Sentence?” In: *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. 2019, pp. 4791–4800.

- [32] Biao Zhang and Rico Sennrich. *Root Mean Square Layer Normalization*. 2019. arXiv: 1910.07467 [cs.LG]. URL: <https://arxiv.org/abs/1910.07467>.